



COMP2129

Week 11 Tutorial

TITLE

Exercise 1: Program Design

Using a pen and paper, explain how you would write the following UNIX utilities: **echo**, **sort**, **uniq**. You should explain it in a high-level way, but also mention any specific C functions which have to be used. You don't have to implement any command-line flags, just the basic functionality. Here is an example of how to explain how to write **cat**:

1. If there are no arguments, just copy the contents of **stdin** to **stdout**, by using an **fread()**, **fwrite()** loop until **fread()** returns **EOF**, then exit.
2. If there are arguments, go through the argument list, trying to open each string as a file in read mode, for example **fopen(argv[1], "r")**. If the file did not open successfully, write an error message to **stderr** using **perror()**. If the file did open successfully, copy its contents to **stdout** using an **fread()** **fwrite()** loop.
3. Return status code 1 if any file could not be opened, and 0 otherwise.

Your tutor will come around during the tutorial and get you to explain one of the utilities, as well as perhaps further implementation details, such as sensible sizes for buffers and how to handle any errors.

Exercise 2: Dining Philosophers 2

The **philosophers.c** file in the **week10_material.zip** archive has an implementation of the dining philosophers in pthreads, with each philosopher modelled as a thread and each fork modelled as a mutex. In the previous tutorial, this was solved by imposing a locking order on the mutexes. This time, solve it by using a global semaphore which only allows a certain number of philosophers to attempt to eat at any one time.

Exercise 3: Diagnosing a deadlock

There are four threads in the following program, labelled T_1 to T_4 , which are (amongst other operations) locking and unlocking three mutexes labelled A , B , and C . Their order of locking is given here:

$T_1 : lock(A), lock(B), lock(C), unlock(A), unlock(B), unlock(C)$

$T_2 : lock(A), lock(B), lock(C), unlock(C), unlock(B), unlock(A)$

$T_3 : lock(B), lock(C), unlock(B), unlock(C)$

$T_4 : lock(C), lock(A), unlock(A), unlock(C)$

Which threads have the potential to deadlock here? Draw two possible interleavings of instructions that cause the threads to go into a deadlock.

Exercise 4: False Sharing

Download the file `falseshare.c` from Piazza or eLearning.

- Describe what the program is doing.
- Identify the parallel and sequential parts of this program.
- Compile and run the program using the `time` utility, for 1, 2, 3, and 4 threads. Does the program speed up as much as it should? Why/why not?
- Identify where false sharing is occurring in this program, and different ways of mitigating it.
- Implement two ways of mitigating false sharing, and use the `time` utility again to ensure your program speeds up as you expect.

Exercise 5: Limits of parallelism

Amdahl's law gives limits on the maximum possible speedup of programs. What is the maximum possible speedup of the following programs:

1. 10% sequential, 90% parallel, with 4 processors used.
2. 90% sequential, 10% parallel, with 32 processors used.
3. 5% sequential, 95% parallel, with an unlimited number of processors.

Discuss some reasons why the maximum theoretical speedup is not often achieved in computations.